

In [3]: `import pandas as pd`

In [4]: `import matplotlib.pyplot as plt`  
`import numpy as np`

In [5]: `df=pd.read_csv("kc_house_data.csv")`

In [6]: `df`

Out[6]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode	lat	
	0	7129300520	20141013T000000	221900.0	3	1.00	1180	5650	1.0	0	0 ...	7	1180.0	0	1955	0	98178	47.5112	-1
	1	6414100192	20141209T000000	538000.0	3	2.25	2570	7242	2.0	0	0 ...	7	2170.0	400	1951	1991	98125	47.7210	-1
	2	5631500400	20150225T000000	180000.0	2	1.00	770	10000	1.0	0	0 ...	6	770.0	0	1933	0	98028	47.7379	-1
	3	2487200875	20141209T000000	604000.0	4	3.00	1960	5000	1.0	0	0 ...	7	1050.0	910	1965	0	98136	47.5208	-1
	4	1954400510	20150218T000000	510000.0	3	2.00	1680	8080	1.0	0	0 ...	8	1680.0	0	1987	0	98074	47.6168	-1
	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
	21608	263000018	20140521T000000	360000.0	3	2.50	1530	1131	3.0	0	0 ...	8	1530.0	0	2009	0	98103	47.6993	-1
	21609	6600060120	20150223T000000	400000.0	4	2.50	2310	5813	2.0	0	0 ...	8	2310.0	0	2014	0	98146	47.5107	-1
	21610	1523300141	20140623T000000	402101.0	2	0.75	1020	1350	2.0	0	0 ...	7	1020.0	0	2009	0	98144	47.5944	-1
	21611	291310100	20150116T000000	400000.0	3	2.50	1600	2388	2.0	0	0 ...	8	1600.0	0	2004	0	98027	47.5345	-1
	21612	1523300157	20141015T000000	325000.0	2	0.75	1020	1076	2.0	0	0 ...	7	1020.0	0	2008	0	98144	47.5941	-1

21613 rows × 21 columns

In [7]: `df.shape`

Out[7]: (21613, 21)

In [38]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     21613 non-null  int64
1   date                   21613 non-null  object
2   price                  21613 non-null  float64
3   bedrooms               21613 non-null  int64
4   bathrooms              21613 non-null  float64
5   sqft_living            21613 non-null  int64
6   sqft_lot               21613 non-null  int64
7   floors                 21613 non-null  float64
8   waterfront             21613 non-null  int64
9   view                   21613 non-null  int64
10  condition              21613 non-null  int64
11  grade                  21613 non-null  int64
12  sqft_above             21611 non-null  float64
13  sqft_basement          21613 non-null  int64
14  yr_built               21613 non-null  int64
15  yr_renovated           21613 non-null  int64
16  zipcode                21613 non-null  int64
17  lat                    21613 non-null  float64
18  long                   21613 non-null  float64
19  sqft_living15          21613 non-null  int64
20  sqft_lot15             21613 non-null  int64
dtypes: float64(6), int64(14), object(1)
memory usage: 3.5+ MB
```

In [48]: `df['date'] = pd.to_datetime(df['date'])`  
`df['Month'] = df['date'].apply(lambda date: date.month)`  
`df['Year'] = df['date'].apply(lambda date: date.year)`

In [107... `Y = df['price'].values`

In [108... `print(X.shape)`  
`print(Y.shape)`  
  
(21611,)
(21611,)

In [109... `df.head()`

Out[109...:

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	...	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	sqft_living15	sqft_lot15
0	7129300520	221900.0	3	1.00	1180	5650	1.0	0	0	3	...	0	1955	0	98178	47.5112	-122.257	1340	5650
1	6414100192	538000.0	3	2.25	2570	7242	2.0	0	0	3	...	400	1951	1991	98125	47.7210	-122.319	1690	7639
2	5631500400	180000.0	2	1.00	770	10000	1.0	0	0	3	...	0	1933	0	98028	47.7379	-122.233	2720	8062
3	2487200875	604000.0	4	3.00	1960	5000	1.0	0	0	5	...	910	1965	0	98136	47.5208	-122.393	1360	5000
4	1954400510	510000.0	3	2.00	1680	8080	1.0	0	0	3	...	0	1987	0	98074	47.6168	-122.045	1800	7503

5 rows × 22 columns

In [110... `df.isnull().sum()`

Out[110...:

id	0
price	0
bedrooms	0

```

bathrooms      0
sqft_living     0
sqft_lot        0
floors          0
waterfront      0
view            0
condition       0
grade           0
sqft_above      0
sqft_basement   0
yr_built        0
yr_renovated     0
zipcode         0
lat             0
long            0
sqft_living15   0
sqft_lot15      0
Month           0
Year            0
dtype: int64

```

```
In [111] df.dropna(inplace=True)
```

```
In [112] df.isnull().sum()
```

```


Out[112] id      0
price      0
bedrooms   0
bathrooms  0
sqft_living 0
sqft_lot   0
floors     0
waterfront 0
view       0
condition  0
grade      0
sqft_above 0
sqft_basement 0
yr_built   0
yr_renovated 0
zipcode    0
lat        0
long       0
sqft_living15 0
sqft_lot15  0
Month      0
Year       0
dtype: int64

```

```
In [113] sqft=df['sqft_living']
```

```
In [114] sqft.plot()
```

```
Out[114] <AxesSubplot:>
```



```

In [167] X = df[['bedrooms','bathrooms','sqft_living','sqft_lot','floors','waterfront','view','condition',
                'grade','sqft_above','sqft_basement','sqft_living15','sqft_lot15']].values
y = df['price'].values

```

```
In [168] X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=50)
```

Normalize the data.

```

In [169] std = StandardScaler()
X = std.fit_transform(X)

```

```
In [170] from sklearn.ensemble import RandomForestRegressor
```

```

In [171] from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score,mean_squared_error,mean_absolute_error
from sklearn.preprocessing import PolynomialFeatures, StandardScaler

```

```

In [172] std = StandardScaler()
X = std.fit_transform(X)

```

```

In [173] rfr = RandomForestRegressor(n_estimators=200)
rfr.fit(X_train,y_train)

```

```
Out[173] RandomForestRegressor(n_estimators=200)
```

```

In [174] score_rfr = rfr.score(X_train,y_train)
prev_rfr = rfr.predict(X_test)
          an_absolute_error(y_test,prev_rfr)

```

```
mse_rfr = mean_squared_error(y_test,prev_rfr)
rmse_rfr = np.sqrt(mean_squared_error(y_test,prev_rfr))
```

```
In [175...
print('Mae: ',mae_rfr)
print('Mse: ',mse_rfr)
print('Rmse: ',rmse_rfr)
```

```
Mae: 123126.961817964
Mse: 37677739823.719635
Rmse: 194107.5470550273
```

## LINEAR REGRESSION

```
In [176...
lr = LinearRegression()
lr.fit(X_train,y_train)
```

```
Out[176... LinearRegression()
```

```
In [177...
pred_lr = lr.predict(X_test)
score_lr = lr.score(X_train,y_train)
```

```
In [178...
mae_lr = mean_absolute_error(y_test,pred_lr)
mse_lr = mean_squared_error(y_test,pred_lr)
rmse_lr = np.sqrt(mse_lr)
```

```
In [179...
print('Mae_lr: ',mae_lr)
print('Mse_lr: ',mse_lr)
print('Rmse_lr: ',rmse_lr)
```


```
Mae_lr: 153081.6455419886
Mse_lr: 54814955728.526276
Rmse_lr: 234125.93988818556
```

```
In [180...
import matplotlib.pyplot as plt
```

```
In [181...
def resizeplot(l,a):
    plt.figure(figsize=(l,a))
```

```
In [195...
resizeplot(10,6)
plt.scatter(y_test,pred_lr)
plt.plot(y_test,y_test,color='red')
```

```
Out[195... [<matplotlib.lines.Line2D at 0x2542230d3c8>]
```



## example of plotting a gradient descent search on a one-dimensional function

```
In [209...
from numpy import asarray
from numpy import arange
from numpy.random import rand
from matplotlib import pyplot

# objective function
def objective(x):
    return x**2.0

# derivative of objective function
def derivative(x):
    return x * 2.0

# gradient descent algorithm
def gradient_descent(objective, derivative, bounds, n_iter, step_size):
    # track all solutions
    solutions, scores = list(), list()
    # generate an initial point
    solution = bounds[:, 0] + rand(len(bounds)) * (bounds[:, 1] - bounds[:, 0])
    # run the gradient descent
    for i in range(n_iter):
        # calculate gradient
        gradient = derivative(solution)
        # take a step
        solution = solution - step_size * gradient
        # evaluate candidate point
        solution_eval = objective(solution)
        # store solution
        solutions.append(solution)
        scores.append(solution_eval)
```

```

        solutions.append(solution)
        scores.append(solution_eval)
        # report progress
        print('>>d f(%s) = %.5f' % (i, solution, solution_eval))
    return [solutions, scores]


# define range for input
bounds = asarray([[-1.0, 1.0]])
# define the total iterations
n_iter = 30
# define the step size
step_size = 0.1
# perform the gradient descent search
solutions, scores = gradient_descent(objective, derivative, bounds, n_iter, step_size)
# sample input range uniformly at 0.1 increments
inputs = arange(bounds[0,0], bounds[0,1]+0.1, 0.1)
# compute targets
results = objective(inputs)
# create a line plot of input vs result
pyplot.plot(inputs, results)
# plot the solutions found
pyplot.plot(solutions, scores, '-.', color='red')
# show the plot
pyplot.show()

```

```

>0 f([-0.45530131]) = 0.20730
>1 f([-0.36424104]) = 0.13267
>2 f([-0.29139284]) = 0.08491
>3 f([-0.23311427]) = 0.05434
>4 f([-0.18649141]) = 0.03478
>5 f([-0.14919313]) = 0.02226
>6 f([-0.11935451]) = 0.01425
>7 f([-0.0954836]) = 0.00912
>8 f([-0.07638688]) = 0.00583
>9 f([-0.06110951]) = 0.00373
>10 f([-0.04888761]) = 0.00239
>11 f([-0.03911008]) = 0.00153
>12 f([-0.03128807]) = 0.00098
>13 f([-0.02503045]) = 0.00063
>14 f([-0.02002436]) = 0.00040
>15 f([-0.01601949]) = 0.00026
>16 f([-0.01281559]) = 0.00016
>17 f([-0.01025247]) = 0.00011
>18 f([-0.00820198]) = 0.00007
>19 f([-0.00656158]) = 0.00004
>20 f([-0.00524927]) = 0.00003
>21 f([-0.00419941]) = 0.00002
>22 f([-0.00335953]) = 0.00001
>23 f([-0.00268762]) = 0.00001
>24 f([-0.0021501]) = 0.00000
>25 f([-0.00172008]) = 0.00000
>26 f([-0.00137606]) = 0.00000
>27 f([-0.00110085]) = 0.00000
>28 f([-0.00088068]) = 0.00000
>29 f([-0.00070454]) = 0.00000

```



In [ ]: